

CreativeHub Claude Task Force

Pre-Workshop Clustering Analysis

Workshop Date: Wednesday | Document: Final, June 2026 | Version: 2 (updated board)

Purpose & Stream Goal

This document summarises every idea submitted to the FigJam board, proposes 7 clusters for Wednesday convergence, and flags the decisions that must be made before or during the workshop. It is also the context document to load into the Claude Advisor project for the scope sprint.

Stream goal: prove Claude is essential to how CreativeHub works and worth keeping on the budget. Every project that survives Wednesday must have a named owner, a real baseline cost, and a measurable KPI.

1. Complete Idea Inventory (24 ideas)

All ideas submitted to the FigJam board by Tuesday. Status: Complete = canvas fully filled; Partial = canvas partially filled; INCOMPLETE = key fields missing; NEW = added in the second board update; PARKING LOT = not ready for Wednesday.

Owner(s)	Idea Title	Type	Status
@Simone Giannotti @Teresa Giannoni	UX Brain / Knowledge OS	Workflow automation	Complete
@Simone Giannotti	Experience Diagnostic Platform	Workflow automation + augmentation	Complete
@AndreaDM	Web Structural Design Rethink	Not filled	PARKING LOT (incomplete)
@Andrea Zito @Francesca Bruschi @Martino Perez	Validated Design Memory	AI design assistant / DesignOps	Complete
@Bayron Vallejo @Giulia Petrolini	Acceleration Framework	Not filled	PARKING LOT (incomplete)
@Federico	Figma Starter Agent	Workflow automation + augmentation	Complete
@AndreaDM	Next-gen AR / Digital Tools	Digital tools	PARKING LOT (incomplete)
@Francesca Bruschi	CRM Modules Generator	Workflow automation + augmentation	Partial
@Francesca Bruschi	Design System Upkeep	Workflow automation + brand guidelines	Complete

@Guillermo de Luna	EDM Copy Review Station	Workflow / quality-control automation	Complete
@AndreaDM @Andrea Zito	Adaptive Commerce Experience	AI-driven personalization layer	Complete — scope flag
@Fabrizio Coco	Brand Memory System	AI-Assisted Creative Workflow	Partial (no KPIs)
@Gaia Sciretti	Concept-to-Visual Mapper	Creative output + workflow augmentation	Complete
@Michele Friolotto	AI Brand Visibility Tracker	Workflow automation + augmentation	Complete (Arnette pilot)
@Gaia Sciretti (@Daniele Catalano credit)	ADHD Reader	Creative output / capability deep-dive	NEW — Complete
@Daniele Catalano @Giulia Petrolini	Face Recognition UI Switcher	Not filled	PARKING LOT (privacy flag)
@Teresa Giannoni @Marco Catani	R&D Insight Autopilot	Workflow / capability	Partial
@Simone Giannotti @Fabio Di Meo	UX Benchmark Sprint	Workflow augmentation	Complete
@Mateo Pinzon @Bayron Vallejo	Brief Estimator	Workflow automation + augmentation	NEW — Complete
@Antonia Helan	Design Timeline Optimizer	Workflow automation / augmentation	Complete
@Marco Catani	Research Augmentation	Capability deep-dive	Complete
@Federico	Smart Asset Retriever	Workflow automation + augmentation	Complete
@Mateo Pinzon @Bayron Vallejo	Design-to-Code Package	Workflow automation + augmentation	Complete
@Andrea Zito @Michela Di Cristina	Events Space Co-pilot	Hybrid: concept generation + decision support	Complete

2. Conflicts to Resolve Before Wednesday

These must be addressed Tuesday evening or at the start of the workshop. They will block team formation if left open.

Ownership overloads

BAYRON VALLEJO is listed on 3 ideas: Acceleration Framework (incomplete), Design-to-Code Package, and Brief Estimator. He must choose one to own. Recommended: merge Brief Estimator into Design-to-Code (Cluster 3) and park Acceleration Framework.

ANDREA ZITO is listed on 3 ideas: Events Space Co-pilot, Validated Design Memory, and Adaptive Commerce Experience. They must choose one. Recommended: Events Space Co-pilot is the most self-contained and has the clearest scope.

GAIA SCIRETTI submitted 2 ideas: ADHD Reader and Concept-to-Visual Mapper. Recommended: merge them into one project (Cluster 5). Both augment creative thinking at different moments: consuming content vs. producing visual output.

ANDREA DM submitted 3+ ideas: Web Structural Design (incomplete), Next-gen AR (incomplete), Adaptive Commerce Experience. Recommended: champion Adaptive Commerce only, and scope it down to one bounded pilot (e.g. personalised product grid for one e-commerce banner). Park the rest.

Ownership swaps (v2 update)

The second board update introduced two ownership changes. Update all cluster materials accordingly.

Idea	Was	Now
Idea	Was	Now
AI Brand Visibility Tracker	@Simone Giannotti @Fabio Di Meo	@Michele Friolotto
UX Benchmark Sprint	@Michele Friolotto	@Simone Giannotti @Fabio Di Meo

3. Proposed Clusters (7 confirmed + 1 optional)

Each cluster groups naturally connected ideas. Rationale is given for every merge. Do not force ideas together where the connection is weak — those are better kept standalone or parked.

CLUSTER 1 DESIGNER'S STARTING POINT

Ideas included	Figma Starter Agent + Smart Asset Retriever
Lead	@Federico
Supporting team	1 UI designer + 1 DesignOps or content person
Merge rationale	Both ideas solve the same moment: before the designer draws anything. One removes friction at file setup, the other at asset retrieval. Same owner, natural merge.
Claude approach	Cowork routine connecting a local Box folder + Figma MCP for file scaffold. Practitioner-demonstrated, not guaranteed end-to-end.
KPI anchor	Time-to-first-design action (combined setup + asset retrieval). Baseline: measure today before Wednesday.

Risk / open question	Box deep-folder integration requires MCP or manual step. Scope the pilot to one campaign folder. Do not claim native real-time Box traversal.
-----------------------------	---

CLUSTER 2 UX KNOWLEDGE & BENCHMARK

Ideas included	UX Brain / Knowledge OS + UX Benchmark Sprint
Lead	@Simone Giannotti @Fabio Di Meo (Benchmark Sprint) + @Teresa Giannoni (Knowledge OS)
Supporting team	Michele Friolotto as methodology contributor (former Benchmark Sprint owner)
Merge rationale	The Knowledge OS is the infrastructure (a shared markdown vault of UX methods, patterns, and decisions). The Benchmark Sprint is the application (Claude retrieves from that vault to generate grounded UI proposals). Phase 1 = build the OS. Phase 2 = run the sprint on top of it.
Claude approach	Claude Projects as the persistent knowledge base. Claude Design for generating UI proposals. Skills for reusable retrieval routines. All officially supported with caveats on Baymard paywall content.
KPI anchor	Time from 'I need a UI reference' to first grounded proposal. Baseline: 30-90 min per task today. Target: under 15 min.
Risk / open question	The Knowledge OS needs a named curator and a minimum viable update cadence before any automation is built on top. If this person is not named Wednesday, the project stalls at Phase 1. Also: Obsidian is local-first and requires a sync layer (Notion, Git, or shared folder) for team-wide access.

CLUSTER 3 DESIGN-TO-DEV GOVERNANCE

Proposed merge: Brief Estimator + Design-to-Code Package + Design System Upkeep form one pipeline: brief arrives > estimate effort > design consistently > hand off to dev. CRM Modules Generator (Francesca) becomes the pilot use case.

Ideas included	Design-to-Code Package + Brief Estimator (both @Mateo @Bayron) + Design System Upkeep (@Francesca). CRM Modules as pilot.
Lead	@Mateo Pinzon (Design-to-Code + pipeline lead) + @Bayron Vallejo (co-lead, must drop Acceleration Framework) + @Francesca Bruschi (DS Upkeep)
Supporting team	Ideally one dev from VisionX as a named stakeholder. Without IT buy-in, the KPIs cannot be measured.
Merge rationale	DS consistency and dev handoff are two sides of the same governance problem. You need the DS in shape to produce a trustworthy handoff package. The Brief Estimator instructs the process before a line is designed. Together they address the core failure: Claude Design is being used to bypass UX, and IT is rebuilding from scratch anyway.
Claude approach	Claude Code for the design-to-code pipeline (officially supported). Claude Projects for the DS knowledge layer. Claude Chat for brief parsing and effort estimation against historical project data (Chat can parse PDFs, Jira exports, Confluence links).

KPI anchor	Accuracy of estimates vs actual delivery time (%). Dev rework reduction (%). % of projects adopting the handoff package. Time saved in scoping (today: approximately 2 weeks per project).
Risk / open question	Bayron has 3 competing ideas. Confirm before Wednesday that he is available for this cluster only. Also: historical project data for the estimator must exist and be accessible. If it does not, the estimator starts with heuristic benchmarks, not real data.

CLUSTER 4 RESEARCH INTELLIGENCE

Ideas included	Research Augmentation + R&D Insight Autopilot
Lead	@Marco Catani (Research Augmentation) + @Teresa Giannoni (R&D Insight Autopilot)
Supporting team	1 additional researcher for coverage
Merge rationale	Both address the research function and both belong to the same team. Research Augmentation = augmenting the researcher's day-to-day (secondary research, methodology challenge, insight-to-deck). R&D Autopilot = connecting the insight repository to external trends. One improves process, the other improves strategic foresight. Complementary, not identical.
Claude approach	Claude Projects with Dovetail exports as the knowledge base. Web search tool for external trend matching. Skills for recurring synthesis templates. All plausible with the caveat that Dovetail connection is not real-time: data must be exported and uploaded manually.
KPI anchor	% of briefs entering creative with a structured secondary-research base. Baseline is plausibly near zero, which makes the delta dramatic and honest. Secondary: time-to-presentable-deck.
Risk / open question	Dovetail export is a manual step. Scope the pilot to one full research cycle (one project from secondary research through insight deck). Do not claim automated Dovetail integration.

CLUSTER 5 BRAND & CREATIVE AUTOMATION

This cluster has two natural sub-tracks. If the team is too large, split it: Gaia leads 'Creative Thinking Augmentation' (ADHD Reader + Concept-to-Visual). Guillermo and Fabrizio run 'Brand Memory & Copy' as a separate project. Both are strong enough standalone.

Ideas included	ADHD Reader (Gaia / Daniele) + Concept-to-Visual Mapper (Gaia) + Brand Memory System (Fabrizio) + EDM Copy Review Station (Guillermo)
Lead options	Sub-track A: @Gaia Sciretti (ADHD Reader + Concept-to-Visual). Sub-track B: @Guillermo de Luna (EDM Review) + @Fabrizio Coco (Brand Memory).
Merge rationale for Gaia's two ideas	ADHD Reader = consume content strategically. Concept-to-Visual Mapper = produce visual directions from abstract concepts. Both augment creative thinking with Claude at different moments. Natural two-phase project: Phase 1 read, Phase 2 produce.

Claude approach	ADHD Reader: Claude Projects + PDF upload + structured prompt for question surfacing (high feasibility). Concept-to-Visual: Claude Design for generating visual explorations from narrative inputs (plausible, not guaranteed quality). Brand Memory: Claude Projects as persistent brand context. EDM Review: Cowork scheduled workflow watching a shared folder, human-in-the-loop review, not autopublish.
KPI anchor	ADHD Reader: % of reading sessions resulting in a captured insight. Concept-to-Visual: time from narrative to first visual direction (target -50 to 70%). EDM Review: review cycle time reduction (target -40 to 50%). Brand Memory: revision rounds per asset (down).
Risk / open question	Brand Memory: Claude Projects holds context per session but does not automatically enforce brand guidelines on all output. Human review stays in the loop. Do not claim 'persistent brand enforcement'. EDM Cowork scheduling requires Claude Desktop to be open on someone's machine at review time.

CLUSTER 6 AI BRAND VISIBILITY TRACKER

Standalone. The most strategically distinctive idea on the board. It is the only one with an existing pilot (Arnette), a clear repeatable methodology, and a problem no one at EL is solving at scale. Protect it.

Idea	AI Brand Visibility Tracker
Lead	@Michele Friolotto (updated ownership from v2 board)
Supporting team	@Fabio Di Meo (methodology contributor) + 1 brand strategist or analyst
What it does	Measures and improves how EL brands are represented across generative AI systems (ChatGPT, Gemini, Perplexity, Google AI, Claude). Generates test queries, scores responses against a fixed framework, maps competitor presence, surfaces prioritized actions. Built on Arnette pilot.
Claude approach	Claude Projects with a structured scoring framework. Web search tool for querying AI platforms in controlled conditions. Claude automates the query generation and scoring; human validates conclusions.
KPI anchor	Hours per brand analysis cycle. Baseline: weeks of manual work per brand. Target: near-replicable execution. Secondary: number of brands analyzed per quarter (scale proof).
Risk / open question	Cross-platform querying at scale raises rate-limit and methodology consistency questions. The Arnette pilot validates the approach but not the automation. Scope the first milestone to automating one full analysis cycle (query generation + scoring) for one additional brand.

CLUSTER 7 PHYSICAL & EVENTS EXPERIENCES

Idea	Events Space Co-pilot
Lead	@Andrea Zito (must drop Adaptive Commerce or Validated Design Memory) + @Michela Di Cristina
Supporting team	1 events/marketing stakeholder as domain expert

What it does	A Claude-powered co-pilot that helps designers create and evaluate exhibition spaces using behavioral insights from past events. Two modes: Concept Booster (generate layout directions from patterns and goals) and Decision Assistant (compare and justify spatial choices). Long-term: data initiative enabling a feedback loop between experience > data > design.
Claude approach	Claude Design for layout concept generation from text descriptions. Claude Projects for evidence-based decision support against historical event data (event reports, movement analytics, stakeholder feedback). Both plausible. The long-term data layer is forward-looking and depends on structured data collection that does not yet exist.
KPI anchor	Time to first layout concept (-30%). Stakeholder alignment score (qualitative, pre/post). Mid-term: dwell time in key areas (+10 to 20%) once data collection is in place.
Risk / open question	The mid-term KPIs depend on structured data (movement, interactions, engagement) that is currently partial or missing. Phase 1 scope must be bounded to 'concepts grounded in past qualitative evidence' rather than 'data-driven layout optimization'. Face Recognition (Daniele/Giulia): still in the parking lot. Unresolved privacy concerns and incomplete canvas. Do not merge into this cluster unless those are addressed.

CLUSTER 8 DESIGN PROCESS & PLANNING (OPTIONAL)

Activate this cluster ONLY if Mateo/Bayron decide NOT to merge the Brief Estimator into Cluster 3. If they do merge it, this cluster does not exist and Antonia joins Cluster 3 instead.

Ideas included	Brief Estimator (@Mateo @Bayron) + Design Timeline Optimizer (@Antonia Helan)
Lead	@Antonia Helan (Timeline Optimizer) + @Mateo Pinzon or @Bayron Vallejo (Estimator)
Merge rationale	The Estimator generates data-informed effort breakdowns from incoming briefs. The Timeline Optimizer stress-tests and defends the resulting plan. Complementary: you use the estimator to set the baseline, then the optimizer to protect and negotiate it.
KPI anchor	Estimation accuracy (%). Reduction in project delays from misestimation. First-time-right delivery rate (projects completed without major redesign loops).
Risk	This project requires historical project data to be accessible and structured. Without it, the estimator uses heuristics only.

4. Parking Lot

These ideas are not ready for Wednesday. They should not receive a pitch slot unless the owner completes the canvas and addresses the noted issues before the session starts.

Idea + Owner	Why parked / what is needed to unpark
Web Structural Design Rethink @AndreaDM	Incomplete: no type, no who/why, no KPI. Could merge into Adaptive Commerce or e-commerce work, but needs a concrete objective first. AndreaDM has too many ideas to run this too.

Acceleration Framework @Bayron @Giulia	Incomplete: no type, no who/why. Bayron is already overloaded on 3 ideas. Park unless Giulia takes full ownership and fills the canvas.
Face Recognition UI Switcher @Daniele @Giulia	Incomplete + unresolved privacy and data ethics concerns (facial recognition of in-store users, biometric data storage). Needs a full privacy architecture section before it can be pitched. Flag for post-workshop R&D exploration if interest remains.
Next-gen AR / Digital Tools @AndreaDM	Incomplete and too technically ambitious for a 6-month task force without a bounded first deliverable. Park as a vision direction. AndreaDM should focus on Adaptive Commerce.
Adaptive Commerce Experience @AndreaDM @Andrea Zito	Complete canvas, but the real-time CRM-to-frontend personalization layer is an engineering architecture project, not a Claude-native one. Claude can help prototype and define the experience, but building the actual orchestration layer is out of scope. Unpark only with a narrow pilot scope: e.g. personalised email content or product grid recommendations for one banner. Also: Andrea Zito must choose between this and Events Co-pilot (Cluster 7).
Experience Diagnostic Platform @Simone Giannotti	Well-defined but ambitious (CX audit platform at EL scale). Simone is the stream lead/facilitator; running a project simultaneously creates a conflict of attention. Options: (a) another team member takes ownership, (b) defer to post-task-force, (c) scope it as a research spike that feeds into Cluster 4.

5. Tuesday Evening Prep Checklist

Lead + board members. ~45 minutes. Board must be ready before 09:00 Wednesday.

- ☐ Move idea stickies into their cluster cards on FigJam. Use cluster numbers and colors from this document.
- ☐ Park incomplete ideas in a separate 'Parking Lot' section on the board with a red dot.
- ☐ Flag the ownership swaps: Michele on Brand Visibility Tracker, Simone/Fabio on UX Benchmark Sprint.
- ☐ Message Bayron directly: confirm he is choosing one idea (Cluster 3) and dropping the others.
- ☐ Message Andrea Zito directly: confirm they are choosing Events Co-pilot (Cluster 7) or another single project.
- ☐ Confirm Gaia is merging ADHD Reader and Concept-to-Visual into one project (Cluster 5).
- ☐ Confirm AndreaDM is focusing on Adaptive Commerce only (and scoping it to a bounded pilot).
- ☐ Pin the scoring rubric on FigJam: (1) Value + KPI clarity, (2) EL relevance, (3) Claude feasibility, (4) Time-to-first-result.
- ☐ Upload this document to the Claude Advisor project knowledge base before Wednesday morning.
- ☐ Print or share the scope sprint starter prompt so each team can open a fresh Advisor chat during the scope sprint block.

6. Using the Claude Advisor in the Scope Sprint

Upload this document to the Claude Advisor project knowledge base. During the scope sprint block (items 5+6 on the agenda), each team opens a fresh chat inside the project and uses this starter prompt:

SCOPE SPRINT STARTER PROMPT (paste into the Advisor): My team has been assigned to [cluster name]. Here is our project idea: [paste your cluster section from this document]. Help us define: 1. A clear one-line objective 2. One measurable KPI with a baseline and target 3. Which Claude surface or workflow makes the most sense to build this 4. Our first testable milestone (what we show, and when) Be honest about what Claude can and cannot do today. Refer to the clustering analysis in the project knowledge base if helpful.

Important: each team should open a fresh chat, not continue a pre-workshop ideation thread. Clean context means the Advisor reads the cluster description clearly without carrying ideation noise into scoping.